

CSED211: Lab. 12

Malloc Lab.

박용곤, 조현욱 TA

csed211-ta@postech.ac.kr

POSTECH

2023.12.11

Table of Contents

- Dynamic Memory Allocation
 - Example: Implicit Free List
- Homework: Malloc Lab
- Quiz
- Practice

What is malloc?

- malloc provide a simple and portable way to **allocate/deallocate a memory block** of desired size
- Linux kernel itself also provides very limited dynamic memory management primitives
 - `brk`, `sbrk` system call
 - Only expand/shrink the end of data segment (just like a stack)
- `libc`, a user-level library, implements malloc using those primitives

libc: malloc & free

- **void *malloc(size_t size)**

- Allocate **size** bytes, and return a pointer to the address of the allocated object
- **my_type *my_obj = (my_type *)malloc(sizeof(my_type)) ;**

- **void free(void *ptr)**

- Free the memory space pointed by **ptr**
- **free(my_obj) ;**

- For more detail, <https://linux.die.net/man/3/malloc>

Challenges in malloc design #1: Execution Speed

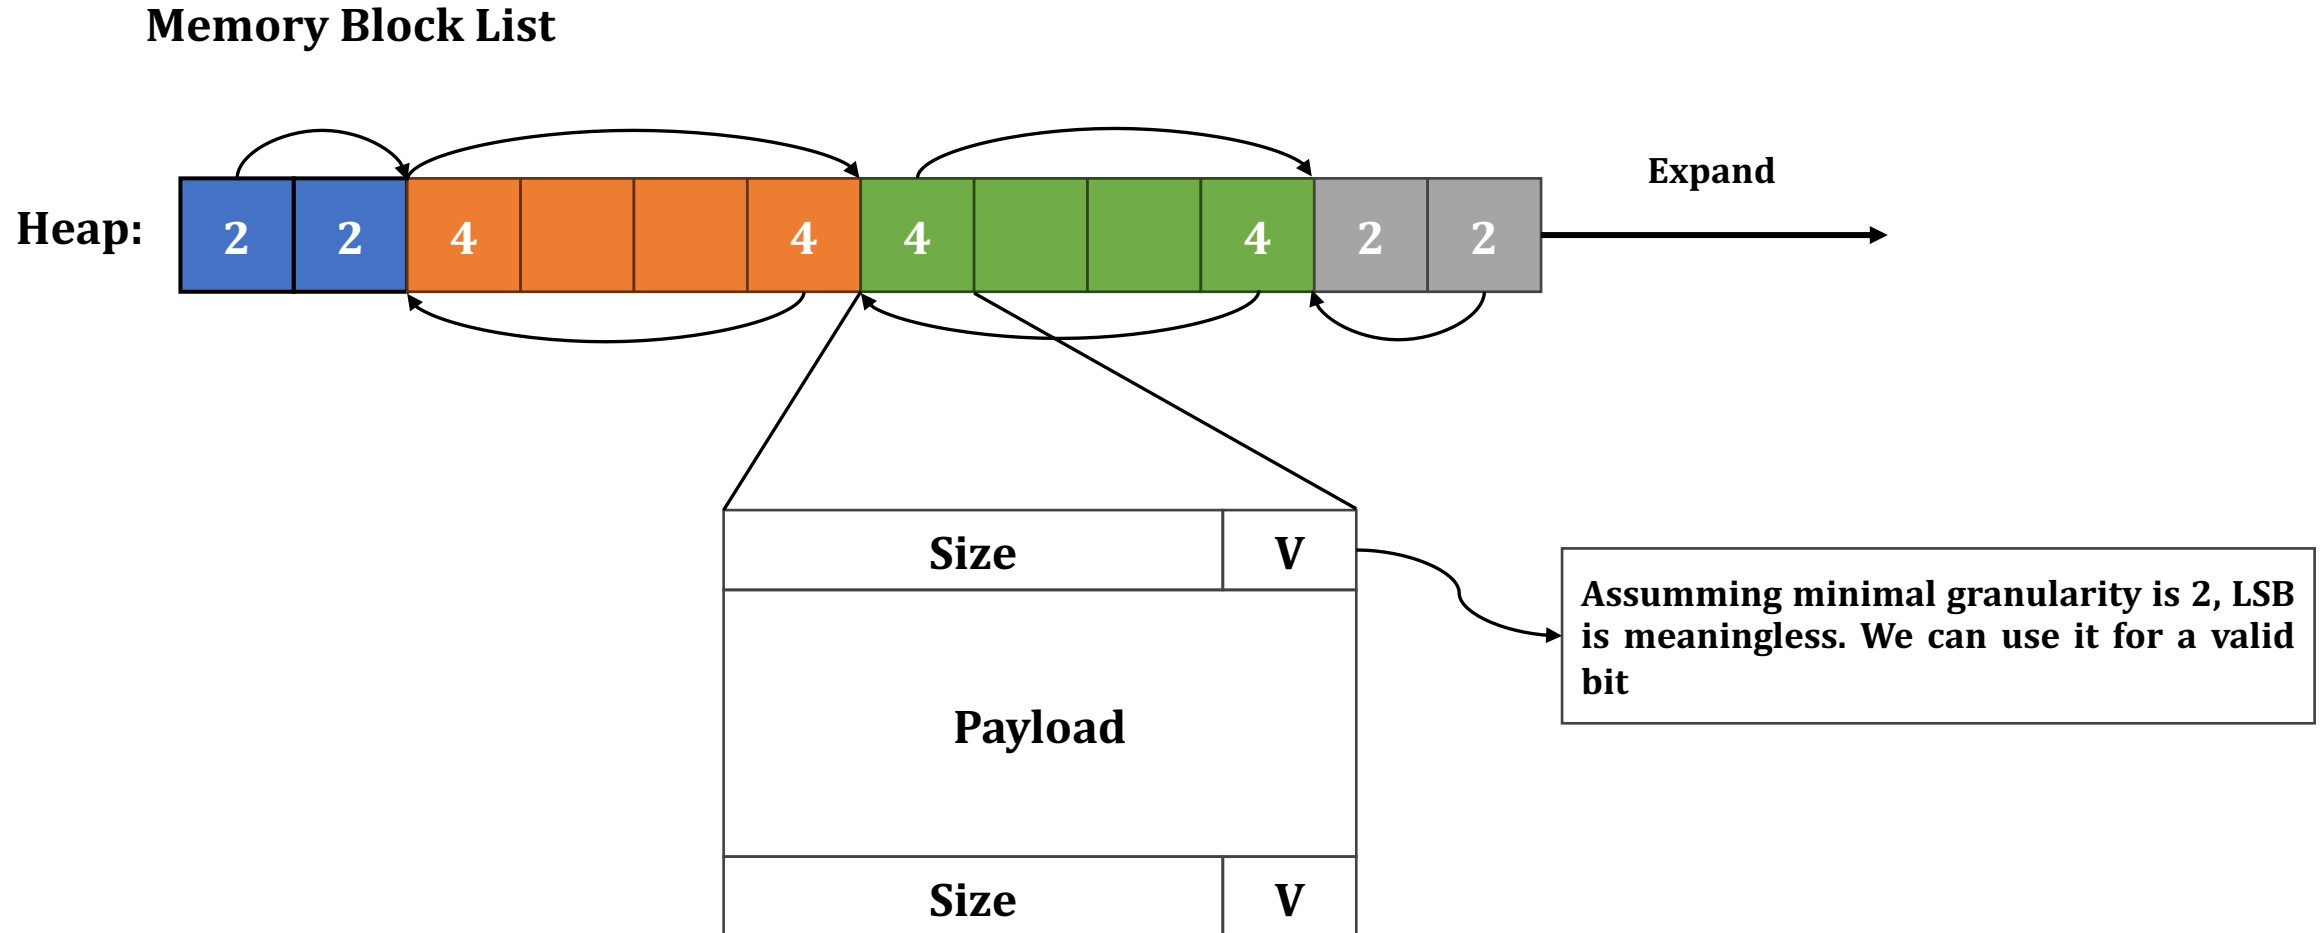
- **How fast** can we:
 - Find a free memory block
 - Release a memory block
- Many kinds of algorithms
 - Bubble sort: $O(n^2)$
 - Merge sort: $O(n \log n)$
 - Linear search: $O(n)$
 - Binary search: $O(\log n)$
 - DFS/BFS: $O(\# \text{ of edges} + \# \text{ of vertices})$

Challenges in malloc design #2: Memory Consumption

- N-byte malloc request consumes at least N-byte
- Data Structure Overhead:
 - Ex) Two words in doubly-linked list (**next** and **prev** pointer)
- Internal Fragmentation:
 - Ex) 3 Byte is requested, but returns 4 Byte (1 Byte wasted)
- External Fragmentation:
 - Ex) There is total 4 Byte of free memory, but increased the heap to satisfy 4 Byte malloc request (4-Byte wasted)


Example: Implicit Free List – Key Idea

- Use **doubly-linked** list to track memory blocks




Example: Implicit Free List - Operations


(1) Initial State




(4-1) free (b) ; Invalidate




(2) d = malloc(2) ; Expand




(4-2) Merge free space (a.k.a. coalesce)



(3) free (c) ; Invalidate



(5) e = malloc(6) ; Split



(You may find the first-fit or the best-fit)

Example: Implicit Free List - Evaluation

- Execution Speed
 - + free = $O(1)$; set invalid & coalesce
 - malloc = $O(\# \text{ of memory blocks})$; linear search
- Memory Space Consumption
 - + Low internal fragmentation
 - + Small data structure overhead (2 words); next & prev displacement
 - + Low external fragmentation if best-fit
 - Severe external fragmentation if first-fit

Practice

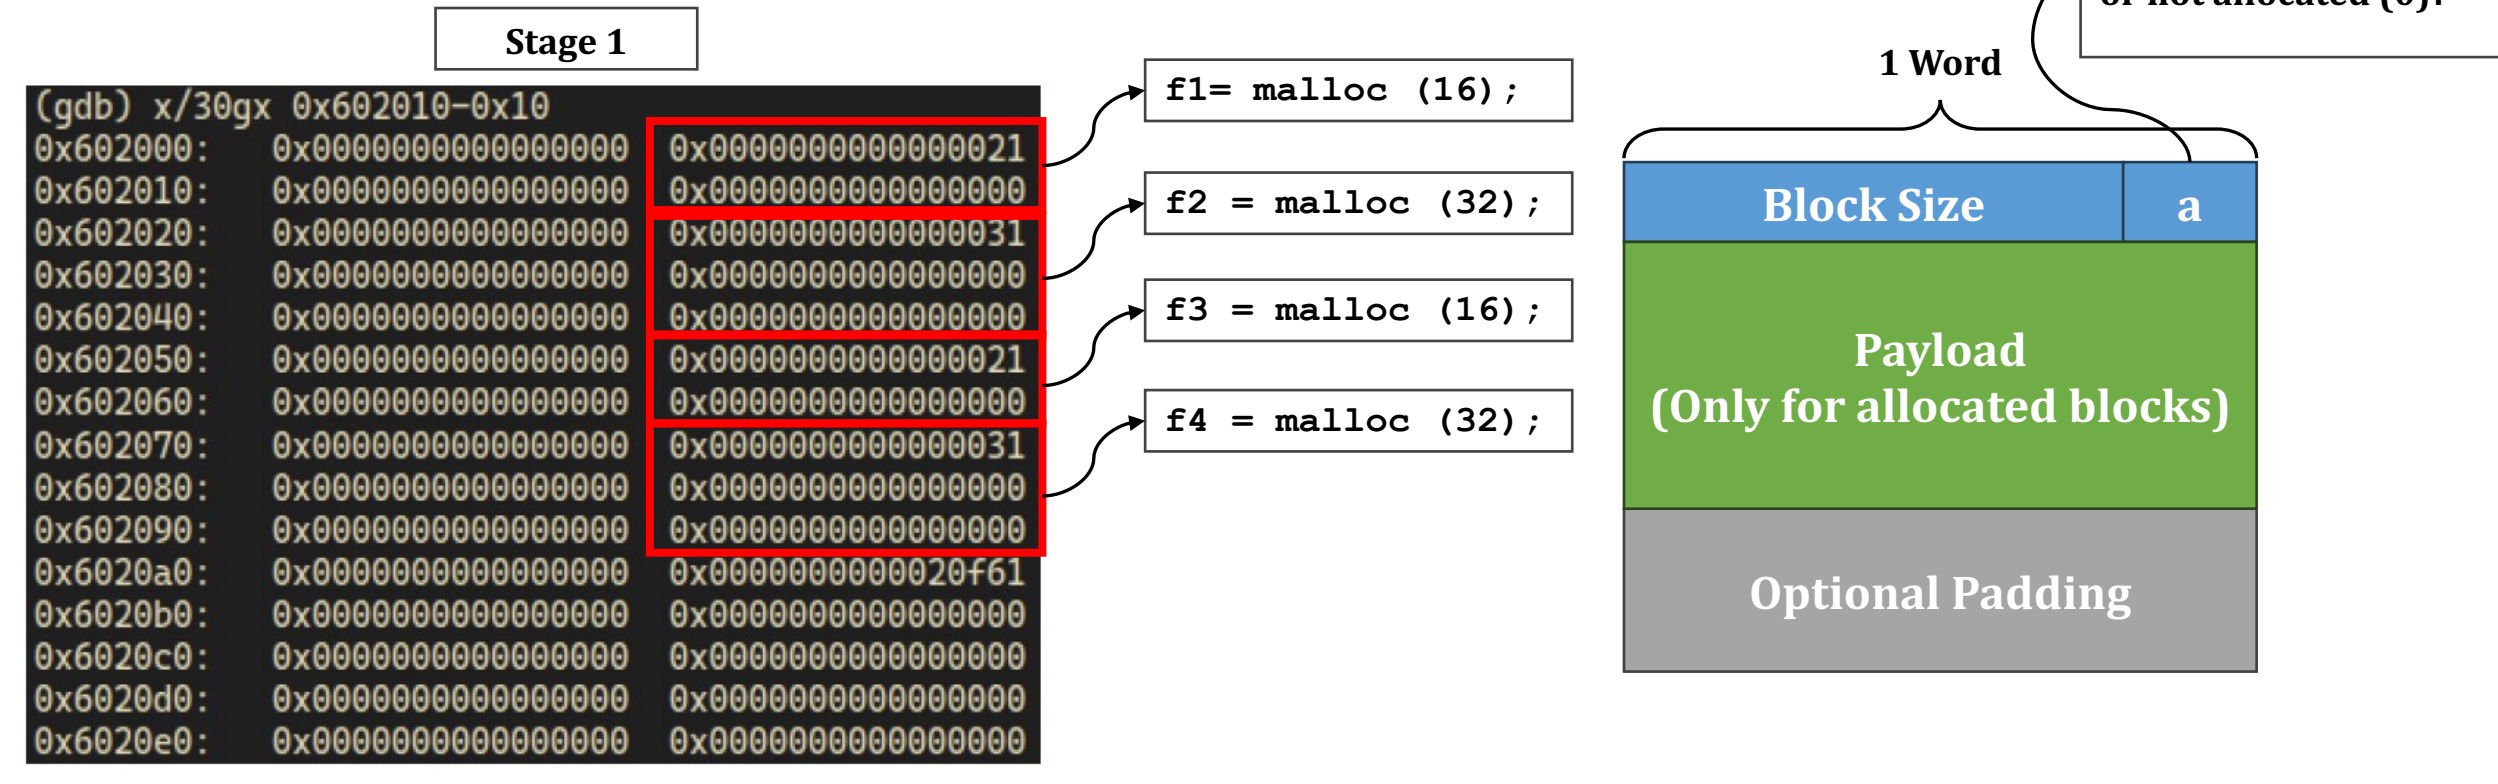
Practice: libc malloc

- We will see how libc malloc works
- Compile and run **practice.c** using gdb
 - `gcc -g practice.c -o practice`
 - `gdb practice`
- Do not try to fully understand this implementation
 - It is very complicated, and you don't need it for assignment

```
1  #include<stdio.h>
2  #include<stdlib.h>
3
4  int main() {
5      void *f1 = malloc(16);
6      printf("0x%x\n", f1);
7      void *f2 = malloc(32);
8      printf("0x%x\n", f2);
9      void *f3 = malloc(16);
10     printf("0x%x\n", f3);
11     void *f4 = malloc(32);
12     printf("0x%x\n", f4);
13     printf("Stage 1\n");
14
15     free(f1);
16     free(f3);
17
18     f1 = malloc(32);
19     printf("0x%x\n", f1);
20
21     free(f2);
22     free(f4);
23
24     return 0;
25 }
26
```

Practice: libc malloc

- libc malloc have a header, and the size is written
- Space overhead for each chunk is 0x10



Quiz

Stage 2

(gdb) x/30gx 0x602010-0x10

a.	<input type="checkbox"/>	0x602000:	0x0000000000000000	0x0000000000000021
		0x602010:	0x0000000000000000	0x0000000000000000
		0x602020:	0x0000000000000000	0x0000000000000031
		0x602030:	0x0000000000000000	0x0000000000000000
		0x602040:	0x0000000000000000	0x0000000000000000
b.	<input type="checkbox"/>	0x602050:	0x0000000000000000	0x0000000000000021
		0x602060:	0x0000000000000000	0x0000000000000000
		0x602070:	0x0000000000000000	0x0000000000000031
		0x602080:	0x0000000000000000	0x0000000000000000
		0x602090:	0x0000000000000000	0x0000000000000000
c.	<input type="checkbox"/>	0x6020a0:	0x0000000000000000	0x00000000000020f61
		0x6020b0:	0x0000000000000000	0x0000000000000000
		0x6020c0:	0x0000000000000000	0x0000000000000000
		0x6020d0:	0x0000000000000000	0x0000000000000000
		0x6020e0:	0x0000000000000000	0x0000000000000000

f1 = malloc(16); free(f1);

f2 = malloc(32);

f3 = malloc(16); free(f3);

f4 = malloc(32);

f1 = malloc(32);
printf("0x%x\n", f1);
→ What is the output?

- a. 0x602000
- b. 0x602050
- c. 0x6020a0
- d. None of the above

Homework

Due: 12/24, 23:59:59

What to do

- You need to complete five functions defined in `mm.c` using **support routines**
 - A detailed explanation is in the writeup file
 - `mm_init`
 - `mm_malloc`
 - `mm_free`
 - `mm_realloc`
 - + `mm_check` (for debugging, **disable when you submit**)
- Complete `mm.c` file and run:
 - `make`
 - `./mdriver` (use `-V` option to get more information)
- Your code will be evaluated in terms of (see 9. Evaluation section in writeup file)
 - Execution speed (**disable `mm_check` for performance, but still `mm_check` will be graded**)
 - Memory space consumption
 - Correctness & style

Grade Policy

- Correctness (20 points)
- Performance (35 points)
 - 5-point deduction for every 10 performance index
 - E.g., 35 points if 100 ~ 91, 30 points if 90 ~ 81, ...
- Report (35 points)
 - Will be graded similarly to Data Lab 2
- Style (10 points) – part of report
 - Heap consistency checker: `mm_check` (5 points)
 - Explain your heap consistency checker in your report with “attached code”
 - Program structure & comments (5 points)
 - Explain your program structure in your report with “attached code”

Results for mm malloc:

trace	valid	util	ops	secs	Kops
0	yes	99%	5694	0.000214	26558
1	yes	98%	5848	0.000216	27087
2	yes	98%	6648	0.000268	24778
3	yes	99%	5380	0.003503	1536
4	yes	99%	14400	0.000259	55556
5	yes	95%	4800	0.000423	11342
6	yes	94%	4800	0.003713	1293
7	yes	55%	12000	0.000268	44726
8	yes	51%	24000	0.002734	8779
9	yes	100%	14401	0.000193	74810
10	yes	87%	14401	0.000229	62941
Total		89%	112372	0.012020	9348

Perf index = 53 (util) + 40 (thru) = 93/100

Grade Policy

- Total 26 traces to grade your malloc
 - Traces are in the /traces directory
 - Each trace will have the same weight

```
[nanimdo@programming2 traces]$ ls
Makefile      cccp.rep      gen_binary2.pl  realloc-bal.rep
README        checktrace.pl gen_coalescing.pl realloc.rep
amptjp-bal.rep coalescing-bal.rep gen_random.pl   realloc2-bal.rep
amptjp.rep    coalescing.rep  gen_realloc.pl  realloc2.rep
binary-bal.rep cp-decl-bal.rep gen_realloc2.pl short1-bal.rep
binary.rep    cp-decl.rep     random-bal.rep  short1.rep
binary2-bal.rep expr-bal.rep    random.rep      short2-bal.rep
binary2.rep   expr.rep        random2-bal.rep short2.rep
cccp-bal.rep  gen_binary.pl   random2.rep
```

Rules (0 point if violated)


- Submission file format
 - Code: [student_id]_mm.c (e.g., 20234567_mm.c)
 - Report: [student_id].pdf (e.g., 20234567.pdf)
 - **DO NOT COMPRESS THE FILES**
- Your code should compile & run successfully
- Return pointers of your allocator must be **8-byte aligned**
- Do not:
 - Change the interface in `mm.c`
 - Use any memory-management related library or system calls
 - Except provided `malloc`, `calloc`, `free`, `realloc`, `sbrk`, `brk`
 - Define any global/static **compound** data structure in `mm.c`
 - E.g., arrays, structs, trees, lists, ...
 - You can define global/static scalar data structure such as `int`, `float`, `pointer`, ...

Q & A

Midterm Exam Statistics

- Personal score will be announced very soon:
 - DO NOT ASK WHEN YOUR SCORE WILL BE ANNOUNCED

Total (300pts)의 히스토그램



Avg.	147.7936
Std.	48.7218136
Min	39.6
Max	279.9